

**A
PROJECT REPORT
ON
PERSONAL LOCAL VOICE
ASSISTANT
USING MACHINE LEARNING**

**A DISSERTATION
SUBMITTED TO THE DEPARTMENT OF COMPUTER SCIENCE AND
ENGINEERING
OF
SILICON INSTITUTE OF TECHNOLOGY, BHUBANESWAR
IN PARTIAL FULFILLMENT OF THE REQUIREMENTS
FOR THE SKILL LAB AND PROJECT**

Dibya Ranjan Rath (SIC-20BCSE62)
Anubhav Mohanty (SIC-20BCSB26)
6th Semester B. Tech.(CSE), Section: B
September 2023

© Copyright by Dibya Ranjan Rath (SIC-20BCSE62)
Anubhav Mohanty (SIC-20BCSB26)
6th Semester B. Tech.(CSE), Section: B 2023
All Rights Reserved

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a Project Report for the 6th Semester B. Tech. SKill Lab and Project.

(Dr. Satyananda Champati Rai) Principal Adviser

Preface

Welcome to "Personal Local Voice Assistant using Machine Learning".

This thesis provides a comprehensive overview of the development of a personal local voice assistant using machine learning. Voice assistants have become increasingly popular in recent years, offering users a convenient and hands-free way to interact with technology. By leveraging machine learning algorithms, a personal local voice assistant can learn and adapt to a user's specific needs and preferences.

In this thesis, we will explore the different machine learning techniques used in the development of a personal local voice assistant, including natural language processing, speech recognition, text-to-speech library and decision tree classifier. We will discuss how these techniques were implemented and optimized, as well as their potential applications in everyday life. We will also evaluate the performance of our voice assistant through user testing.

We hope that this thesis will serve as a valuable resource for those interested in developing personal local voice assistants and advancing the field of machine learning in natural language processing and speech recognition.

Acknowledgments

We would like to express our sincere gratitude to all those who have supported us during the course of this project.

Firstly, we would like to thank our project guide, Dr. Satyananda Champati Rai, for his invaluable guidance, support, and encouragement throughout the project. His expertise and insight have been invaluable in shaping this thesis into its final form.

We would also like to thank the Lab Attendant, Mr. Dibakar Pradhan, for his support and guidance, who provided us with insightful feedback and suggestions throughout the project.

We would like to express our heartfelt thanks to our family and friends for their unwavering support, encouragement, and patience throughout this project. Their constant support and belief in us have been a source of inspiration and motivation.

Finally, we would like to thank all our users who participated in the testing of our voice assistant and provided us with valuable feedback.

Once again, thank you to all those who have contributed to the successful completion of this project. We hope that our personal local voice assistant can be of great use to people and make their everyday life easier.

Contents

Preface	iv
Acknowledgments	v
1 Introduction to Machine Learning	1
1.1 Supervised Learning	1
1.2 Un-Supervised Learning	1
1.3 Classification	2
1.4 Regression	2
1.5 Error/Analysis	2
1.6 Reinforcement Learning	2
1.7 Objectives	2
1.8 Chapter Outline	3
1.9 Conclusion	3
2 Mathematical foundation of ML	4
2.1 Matrix Calculus	4
2.2 Eigen values & Eigen vector	5
2.3 Random variable and distribution	5
2.3.1 Normal	6
2.3.2 Gaussian	7
2.4 Kernel	8
2.4.1 Polynomial Kernel	8
2.4.2 Gaussian Kernel	8
3 Implementation of Machine Learning Algorithms	9
3.1 K-Nearest Neighbour - Classification	9
3.1.1 Pseudocode	9
3.1.2 Code	9

3.2	K-Nearest Neighbour - Regression	10
3.2.1	Pseudocode	10
3.2.2	Code	11
3.3	Multiple Linear Regression	12
3.3.1	Pseudocode	12
3.3.2	Code	12
3.4	Ridge Regression	13
3.4.1	Pseudocode	13
3.4.2	Code	13
3.5	LASSO Regression	14
3.5.1	Pseudocode	14
3.6	Logistic Regression	15
3.6.1	Pseudocode	16
3.6.2	Code	16
3.7	Support Vector Machine - Classification	17
3.7.1	Pseudocode	17
3.8	Support Vector Machine - Regression	19
3.8.1	Pseudocode	19
3.9	Decision Tree for Classification	20
3.9.1	Pseudocode	21
3.9.2	Code	21
3.10	Decision Tree for Regression	22
3.10.1	Pseudocode	22
3.10.2	Code	22
3.11	Naive Bayes' Classifier	23
3.11.1	Pseudocode	23
3.11.2	Code	24
3.12	Random Forest	25
3.12.1	Pseudocode	25
3.12.2	Code	25
3.13	K-Means Clustering	26
3.13.1	Pseudocode	26
3.13.2	Code	27
3.14	Spectral Clustering	28
3.14.1	Pseudocode	28
3.14.2	Code	29
3.15	MDP: Markov Decision Process	29

3.15.1 Pseudocode	29
3.16 MRP: Material Requirement Planning	30
3.16.1 Pseudocode	30
3.17 Queue-learning	31
3.17.1 Pseudocode	31
3.18 SARSA Learning	31
3.18.1 Pseudocode	32
4 Proposed Model	33
4.1 Background	33
4.2 Limitation of Research Article	34
4.3 Design of model	34
4.3.1 Pseudocode for Decision Tree Classifier	34
4.3.2 pseudocode for Personal Local Voice Assistant	35
4.4 Program	36
4.5 Data set Selected	39
4.6 Performance Analysis	39
4.7 Conclusion	42
5 Future Work and conclusion	43
5.1 Future Work	43
5.2 Conclusion	43

List of Tables

List of Figures

4.1	Decision Tree Classifier	40
4.2	Sample Output1	40
4.3	Sample Output2	41
4.4	Sample Output3	41
4.5	Sample Output4	41

Chapter 1

Introduction to Machine Learning

Machine learning (ML) is a branch of artificial intelligence (AI) that enables computers to “self-learn” from training data and improve over time, without being explicitly programmed. Machine learning algorithms are able to detect patterns in data and learn from them, in order to make their own predictions.

1.1 Supervised Learning

Supervised learning, also known as supervised machine learning, is a subcategory of machine learning and artificial intelligence. It is defined by its use of labeled datasets to train algorithms that to classify data or predict outcomes accurately. As input data is fed into the model, it adjusts its weights until the model has been fitted appropriately, which occurs as part of the cross validation process. Supervised learning helps organizations solve for a variety of real-world problems at scale, such as classifying spam in a separate folder from your inbox.

1.2 Un-Supervised Learning

Unsupervised learning, also known as unsupervised machine learning, uses machine learning algorithms to analyze and cluster unlabeled data sets. These algorithms discover hidden patterns or data groupings without the need for human intervention. Its ability to discover similarities and differences in information make it the ideal solution for exploratory data analysis, cross-selling strategies, customer segmentation, and image recognition.

1.3 Classification

The Classification algorithm is a Supervised Learning technique that is used to identify the category of new observations on the basis of training data. In Classification, a program learns from the given data set or observations and then classifies new observation into a number of classes or groups. Such as, Yes or No, 0 or 1, Spam or Not Spam, cat or dog, etc. Classes can be called as targets/labels or categories.

1.4 Regression

Regression analysis is a statistical method to model the relationship between a dependent (target) and independent (predictor) variables with one or more independent variables. More specifically, Regression analysis helps us to understand how the value of the dependent variable is changing corresponding to an independent variable when other independent variables are held fixed. It predicts continuous/real values such as temperature, age, salary, price, etc.

1.5 Error/Analysis

Error analysis is the process of isolating, observing, and diagnosing erroneous ML predictions. The ideal result is that we're able to better understand pockets of high and low performance in the model. When it is stated that "The model accuracy is 90%", it might not be uniform across subgroups of data, and some input conditions could negatively impact the model. Hence, the next step from focusing on aggregate metrics is to review model errors in more detail for improvement.

1.6 Reinforcement Learning

Reinforcement Learning (RL) is the science of decision making. It is about learning the optimal behavior in an environment to obtain maximum reward. In RL, the data is accumulated from machine learning systems that use a trial-and-error method. Data is not part of the input that we would find in supervised or unsupervised machine learning.

1.7 Objectives

The goal of machine learning is often — though not always — to train a model on historical, labelled data (i.e., data for which the outcome is known) in order to predict the value of some quantity on the basis of a new data item for which the target value or classification is unknown. We might, for example, want to predict the lifetime value of customer XYZ, or to predict whether a transaction is fraudulent or not.

1.8 Chapter Outline

The thesis on Personal Local Voice Assistant using Machine Learning aims to develop a personalized and intelligent voice assistant using machine learning algorithms. It begins with an introduction to the importance of voice assistants in everyday life and a review of existing voice assistant technologies. The thesis then provides a literature review of machine learning techniques used in the development of voice assistants, including natural language processing, speech recognition, and decision tree classifier. The data requirements for the development of a personalized voice assistant are discussed, along with data preprocessing techniques and data visualization and exploration. Machine learning models are then presented, including the use of natural language processing and speech recognition for voice assistant tasks. The evaluation metrics for assessing model performance and the experimental design and methodology are also discussed. The experimental results and performance evaluation of the machine learning models are presented, including comparisons with existing voice assistant technologies. The thesis concludes with a summary of the main findings, implications and limitations of the study, recommendations for future research, and technical information in the appendix.

1.9 Conclusion

Machine Learning is directly or indirectly involved in our daily routine. We have seen various machine learning applications that are very useful for surviving in this technical world. Although machine learning is in the developing phase, it is continuously evolving rapidly. The best thing about machine learning is its High-value predictions that can guide better decisions and smart actions in real-time without human intervention. Hence, at the end of this article, we can say that the machine learning field is very vast, and its importance is not limited to a specific industry or sector; it is applicable everywhere for analyzing or predicting future events.

Chapter 2

Mathematical foundation of ML

In this Chapter we are going to learn about various Mathematical foundation of ML such as:-

2.1 Matrix Calculus

Matrix calculus is a specialized notation for doing multi-variable calculus, especially over spaces of matrices. It collects the various partial derivatives of a single function with respect to many variables, and/or of a multivariate function with respect to a single variable, into vectors and matrices that can be treated as single entities. Suppose we have two matrices \mathbf{A} and \mathbf{B} , and a scalar c . Here are some examples of matrix calculus operations:

1. Matrix addition: $\mathbf{A} + \mathbf{B} = \begin{pmatrix} a_{11} + b_{11} & a_{12} + b_{12} \\ a_{21} + b_{21} & a_{22} + b_{22} \end{pmatrix}$
2. Scalar multiplication: $c\mathbf{A} = \begin{pmatrix} ca_{11} & ca_{12} \\ ca_{21} & ca_{22} \end{pmatrix}$
3. Matrix multiplication: $\mathbf{AB} = \begin{pmatrix} a_{11}b_{11} + a_{12}b_{21} & a_{11}b_{12} + a_{12}b_{22} \\ a_{21}b_{11} + a_{22}b_{21} & a_{21}b_{12} + a_{22}b_{22} \end{pmatrix}$
4. Transpose: $\mathbf{A}^\top = \begin{pmatrix} a_{11} & a_{21} \\ a_{12} & a_{22} \end{pmatrix}$
5. Trace: $\text{tr}(\mathbf{A}) = a_{11} + a_{22}$
6. Determinant: $\det(\mathbf{A}) = a_{11}a_{22} - a_{12}a_{21}$
7. Inverse: $\mathbf{A}^{-1} = \frac{1}{a_{11}a_{22} - a_{12}a_{21}} \begin{pmatrix} a_{22} & -a_{12} \\ -a_{21} & a_{11} \end{pmatrix}$

2.2 Eigen values & Eigen vector

An eigenvector or characteristic vector of a linear transformation is a nonzero vector that changes at most by a scalar factor when that linear transformation is applied to it. The corresponding eigenvalue is the factor by which the eigenvector is scaled. Suppose we have a matrix \mathbf{A} of size $n \times n$. An eigenvector \mathbf{v} and corresponding eigenvalue λ satisfy the following equation:

$$\mathbf{A}\mathbf{v} = \lambda\mathbf{v}$$

To find the eigenvalues and eigenvectors of \mathbf{A} , we first solve the characteristic equation:

$$\det(\mathbf{A} - \lambda\mathbf{I}) = 0$$

where \mathbf{I} is the identity matrix of size $n \times n$. The solutions to this equation are the eigenvalues of \mathbf{A} .

Next, we find the eigenvectors associated with each eigenvalue by solving the equation:

$$(\mathbf{A} - \lambda\mathbf{I})\mathbf{v} = \mathbf{0}$$

where $\mathbf{0}$ is the zero vector. The solutions to this equation are the eigenvectors associated with the eigenvalue λ .

The eigenvalues and eigenvectors of \mathbf{A} are then given by the set of pairs $\{(\lambda_i, \mathbf{v}_i)\}$, where λ_i is the i -th eigenvalue and \mathbf{v}_i is the corresponding eigenvector.

2.3 Random variable and distribution

A random variable is a numerical description of the outcome of a statistical experiment. The probability distribution for a random variable describes how the probabilities are distributed over the values of the random variable.

Suppose we have a random variable X that takes on values from a sample space Ω . The probability distribution of X is a function that assigns probabilities to each of the possible values of X . This function is called the probability mass function (PMF) for a discrete random variable, or the probability density function (PDF) for a continuous random variable.

For a discrete random variable, the PMF is defined as:

$$p_X(x_i) = P(X = x_i)$$

where x_i is a possible value of X . The PMF must satisfy the following conditions:

1. $0 \leq p_X(x_i) \leq 1$ for all i

$$2. \sum_i p_X(x_i) = 1$$

For a continuous random variable, the PDF is defined as:

$$f_X(x) \geq 0 \quad \text{and} \quad \int_{-\infty}^{\infty} f_X(x) dx = 1$$

The probability that X takes on a value in the interval $[a, b]$ is given by:

$$P(a \leq X \leq b) = \int_a^b f_X(x) dx$$

Some common probability distributions include:

- Bernoulli distribution: a discrete distribution with two possible outcomes, often used to model success/failure experiments
- Binomial distribution: a discrete distribution with n independent Bernoulli trials, often used to model the number of successes in a fixed number of trials
- Normal distribution: a continuous distribution with a bell-shaped curve, often used to model naturally occurring phenomena

2.3.1 Normal


Normal distribution is a bell-shaped curve, and it is assumed that during any measurement values will follow a normal distribution with an equal number of measurements above and below the mean value. The probability density function of the normal distribution, also known as the Gaussian distribution, is given by:

The probability density function of the normal distribution is given by:

The normal distribution, also known as the Gaussian distribution, is a probability distribution that is often used in statistical analysis. It is characterized by two parameters: the mean μ and the standard deviation σ . The probability density function of the normal distribution is given by:

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{1}{2} \left(\frac{x - \mu}{\sigma}\right)^2\right)$$

where x is the random variable, μ is the mean, and σ is the standard deviation.



The normal distribution is symmetric around the mean, with the highest probability density occurring at the mean. The standard deviation determines the spread of the distribution, with larger values resulting in a wider distribution.

2.3.2 Gaussian

Gaussian distribution is a bell-shaped curve, and it is assumed that during any measurement values will follow a normal distribution with an equal number of measurements above and below the mean value. The Gaussian distribution, also known as the normal distribution, is a probability distribution that is often used in statistics and machine learning. It has the following probability density function:

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

where μ is the mean of the distribution and σ^2 is its variance.

We can also express the Gaussian distribution in terms of its cumulative distribution function (CDF), which gives the probability that a random variable X drawn from the distribution is less than or equal to a given value x :

$$F(x) = \frac{1}{2} \left[1 + \operatorname{erf} \left(\frac{x-\mu}{\sqrt{2}\sigma} \right) \right]$$

where $\operatorname{erf}(x)$ is the error function.

The mean and variance of the Gaussian distribution are given by:

$$\begin{aligned} \mathbb{E}[X] &= \mu \\ \operatorname{Var}[X] &= \sigma^2 \end{aligned}$$

The Gaussian distribution is often used as a model for many natural phenomena, such as the distribution of heights and weights in a population, the errors in scientific measurements, and the noise in electronic signals.

2.4 Kernel

Kernel is a computer program at the core of a computer's operating system and generally has complete control over everything in the system. It is the portion of the operating system code that is always resident in memory and facilitates interactions between hardware and software components.

2.4.1 Polynomial Kernel

The polynomial kernel is a kernel function commonly used with support vector machines and other kernelized models, that represents the similarity of vectors in a feature space over polynomials of the original variables, allowing learning of non-linear models. The polynomial kernel is a type of kernel function used in machine learning for support vector machines (SVMs) and other kernel-based methods. It has the form:

$$K(x, y) = (\langle x, y \rangle + c)^d$$

where x and y are feature vectors, $\langle x, y \rangle$ denotes the dot product of the vectors, c is a constant, and d is the degree of the polynomial.

In SVMs, the polynomial kernel is often used to handle non-linearly separable data by transforming the data into a higher-dimensional feature space, where it may be linearly separable. The degree of the polynomial determines the degree of the transformation, and the constant c controls the influence of low-degree versus high-degree polynomial terms.

2.4.2 Gaussian Kernel

The Gaussian kernel is the physical equivalent of the mathematical point. It is not strictly local, like the mathematical point, but semi-local. It has a Gaussian weighted extent, indicated by its inner scale s . The Gaussian kernel is a popular choice for smoothing data and estimating probability density functions. It is defined as:

$$K(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right)$$

where x is the distance from the center of the kernel, and σ is a parameter that controls the width of the kernel. Larger values of σ result in a wider kernel that smooths over a larger range of values.

The Gaussian kernel can be used in a variety of applications, such as image processing, machine learning, and signal processing. In machine learning, it is commonly used as a similarity measure in kernel methods, such as support vector machines and Gaussian processes.

Chapter 3

Implementation of Machine Learning Algorithms

3.1 K-Nearest Neighbour - Classification

K-Nearest Neighbour classification is a supervised learning algorithm used for classification and regression analysis. It finds the k-nearest neighbors of new input data based on a distance metric and assigns it to the class that is most common among the k neighbors. It is simple to implement but can be computationally expensive and its performance depends on the choice of k and distance metric.

3.1.1 Pseudocode

```
1: procedure KNNCLASSIFICATION( $X, Y, x_{test}, k$ )
2:    $D \leftarrow$  array of size  $n$ , where  $n$  is the number of training examples
3:   for  $i \leftarrow 1$  to  $n$  do
4:      $D_i \leftarrow$  distance between  $x_{test}$  and  $X_i$ 
5:   end for
6:    $indices \leftarrow$  indices of the  $k$  smallest values in  $D$ 
7:    $Y_{neighbors} \leftarrow$  classes of the  $k$  nearest neighbors, based on their indices in  $Y$ 
8:    $y_{pred} \leftarrow$  class with the highest frequency in  $Y_{neighbors}$ 
9:   return  $y_{pred}$ 
10: end procedure
```

3.1.2 Code

```
import pandas as pd
```

```

from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

Load the dataset
df = pd.read_csv('housing.csv')

Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

Train the KNN model
k = 5
knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(X_train, y_train)

Predict the classes of the testing set
y_pred = knn.predict(X_test)

Evaluate the performance of the model
accuracy = accuracy_score(y_test, y_pred)
print('Accuracy:', accuracy)

Accuracy: 0.01020408163265306

```

3.2 K-Nearest Neighbour - Regression

KNN regression is a non-parametric method that, in an intuitive manner, approximates the association between independent variables and the continuous outcome by averaging the observations in the same neighbourhood.

3.2.1 Pseudocode

-
- 1: **procedure** KNNREGRESSION(X, Y, x_{test}, k)
 - 2: $D \leftarrow$ array of size n , where n is the number of training examples
 - 3: **for** $i \leftarrow 1$ to n **do**

```

4:      $D_i \leftarrow$  distance between  $x_{test}$  and  $X_i$ 
5: end for
6:      $indices \leftarrow$  indices of the  $k$  smallest values in  $D$ 
7:      $Y_{neighbors} \leftarrow$  values of the  $k$  nearest neighbors, based on their indices in  $Y$ 
8:      $y_{pred} \leftarrow$  mean of  $Y_{neighbors}$ 
9:     return  $y_{pred}$ 
10: end procedure

```

3.2.2 Code

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error

# Load the dataset
df = pd.read_csv('housing.csv')

# Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Train the KNN model
k = 5
knn = KNeighborsRegressor(n_neighbors=k)
knn.fit(X_train, y_train)

# Predict the target values of the testing set
y_pred = knn.predict(X_test)

# Evaluate the performance of the model
mse = mean_squared_error(y_test, y_pred)
print('Mean squared error:', mse)

```

Mean squared error: 4014687600.0

3.3 Multiple Linear Regression

Multiple linear regression (MLR), also known simply as multiple regression, is a statistical technique that uses several explanatory variables to predict the outcome of a response variable. Multiple regression is an extension of linear (OLS) regression that uses just one explanatory variable.

3.3.1 Pseudocode

```

1: procedure MLR( $X, Y$ )
2:    $n, p \leftarrow$  dimensions of  $X$ 
3:    $\bar{X} \leftarrow$  mean of each column in  $X$ 
4:    $\bar{Y} \leftarrow$  mean of  $Y$ 
5:    $S_{xx} \leftarrow \sum_{i=1}^n (X_i - \bar{X})(X_i - \bar{X})^T$ 
6:    $S_{xy} \leftarrow \sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})$ 
7:    $B \leftarrow (S_{xx})^{-1} S_{xy}$ 
8:    $b_0 \leftarrow \bar{Y} - B^T \bar{X}$ 
9:   return  $b_0, B$ 
10: end procedure

```

3.3.2 Code

```

import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score

# Load the dataset
df = pd.read_csv('housing.csv')

# Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Train the MLR model
mlr = LinearRegression()
mlr.fit(X_train, y_train)

```

```

# Predict the target values of the testing set
y_pred = mlr.predict(X_test)

# Evaluate the performance of the model
r2 = r2_score(y_test, y_pred)
print('R^2:', r2)

```

```
R^2: 0.6910934003098523
```

3.4 Ridge Regression

Ridge Regression is a type of linear regression used to prevent overfitting by adding a penalty term to the sum of squared residuals. This penalty term is the L2 norm of the coefficients multiplied by a hyper parameter lambda. It results in a more generalized model that is less likely to overfit the training data, making it useful for high-dimensional datasets with multicollinearity.

3.4.1 Pseudocode

```

1: procedure RIDGEGRESSION( $X, Y, \lambda$ )
2:    $n \leftarrow$  number of training examples
3:    $m \leftarrow$  number of features in  $X$ 
4:    $I \leftarrow$  identity matrix of size  $m \times m$ 
5:    $w \leftarrow (X^T X + \lambda I)^{-1} X^T Y$ 
6:   return  $w$ 
7: end procedure

```

3.4.2 Code

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error

# Load the dataset
df = pd.read_csv('housing.csv')

# Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Train the Ridge Regression model
ridge = Ridge(alpha=1)
ridge.fit(X_train, y_train)

# Predict the target values of the testing set
y_pred = ridge.predict(X_test)

# Evaluate the performance of the model
mse = mean_squared_error(y_test, y_pred)
print('Mean Squared Error:', mse)

```

Mean Squared Error: 6787779698.377267

3.5 LASSO Regression

Lasso regression is a regularization technique. It is used over regression methods for a more accurate prediction. This model uses shrinkage. Shrinkage is where data values are shrunk towards a central point as the mean. The lasso procedure encourages simple, sparse models (i.e. models with fewer parameters).

3.5.1 Pseudocode

```

1: procedure LASSO( $X, y, \lambda, \alpha, max\_iter$ )
2:    $n, p \leftarrow$  dimensions of  $X$ 
3:    $w \leftarrow$  array of size  $p$  initialized with zeros
4:    $intercept \leftarrow$  mean of  $y$ 
5:    $X \leftarrow$  normalized version of  $X$ 
6:    $X^T \leftarrow$  transpose of  $X$ 
7:    $t \leftarrow 0$ 
8:   while  $t < max\_iter$  do
9:      $t \leftarrow t + 1$ 
10:    for  $j \leftarrow 1$  to  $p$  do
11:       $X_j \leftarrow$  column  $j$  of  $X$ 
12:       $R_j \leftarrow y - intercept - X^T \cdot w + w_j \cdot X_j$ 
13:       $z_j \leftarrow X_j^T \cdot R_j / n + w_j$ 

```



```

14:          $w_j \leftarrow \text{soft threshold function}(z_j, \alpha \cdot \lambda)$ 
15:     end for
16: end while
17: return  $w, \text{intercept}$ 
18: end procedure

```

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Lasso
from sklearn.metrics import mean_squared_error

# Load the dataset
df = pd.read_csv('housing.csv')

# Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Train the Lasso Regression model
lasso = Lasso(alpha=1)
lasso.fit(X_train, y_train)

# Predict the target values of the testing set
y_pred = lasso.predict(X_test)

# Evaluate the performance of the model
mse = mean_squared_error(y_test, y_pred)
print('Mean Squared Error:', mse)

```

Mean Squared Error: 6789024186.532051

3.6 Logistic Regression

Logistic regression is an example of supervised learning. It is used to calculate or predict the probability of a binary (yes/no) event occurring.

3.6.1 Pseudocode

```

1: procedure LOGISTICREGRESSION( $X, Y, \alpha, \epsilon$ )
2:    $\theta \leftarrow$  array of size  $m$ , where  $m$  is the number of features in  $X$ 
3:    $cost \leftarrow$  initial cost value
4:   while  $cost > \epsilon$  do
5:      $h \leftarrow$  sigmoid function applied to  $\theta^T X$ 
6:      $J \leftarrow$  cost function of logistic regression
7:      $\theta \leftarrow \theta - \alpha \frac{\partial J}{\partial \theta}$ 
8:      $cost \leftarrow$  new cost value
9:   end while
10:  return  $\theta$ 
11: end procedure

```

3.6.2 Code

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Load the dataset
df = pd.read_csv('housing.csv')

# Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Train the Logistic Regression model
logreg = LogisticRegression()
logreg.fit(X_train, y_train)

# Predict the target values of the testing set
y_pred = logreg.predict(X_test)

# Evaluate the performance of the model
accuracy = accuracy_score(y_test, y_pred)

```

```
print('Accuracy:', accuracy)
```

```
Accuracy: 0.02040816326530612
```

3.7 Support Vector Machine - Classification

Support Vector Machine (SVM) is a popular machine learning algorithm used for classification, regression, and outlier detection. The main idea of SVM is to find the optimal hyperplane that separates two classes in a high-dimensional space. SVM seeks to maximize the margin between the hyperplane and the closest data points from each class, which are called support vectors. The optimal hyperplane is the one that maximizes the margin while correctly classifying the training data.

In classification problems, SVM can handle both linearly separable and non-linearly separable data by using a kernel function to map the data to a higher-dimensional space where it is more likely to be linearly separable. Common kernel functions include linear, polynomial, and radial basis function (RBF) kernels.

3.7.1 Pseudocode

```

1: procedure SVMCLASSIFICATION( $X, Y, C, kernel, x_{test}$ )
2:    $n \leftarrow$  number of training examples
3:    $m \leftarrow$  number of features
4:    $\alpha \leftarrow$  array of size  $n$ , initialized to 0
5:    $b \leftarrow 0$ 
6:    $E \leftarrow$  array of size  $n$ , initialized to 0
7:    $K \leftarrow$  matrix of size  $n \times n$ , where  $K_{i,j} = kernel(X_i, X_j)$ 
8:   repeat
9:     for  $i \leftarrow 1$  to  $n$  do
10:       $E_i \leftarrow (\sum_{j=1}^n \alpha_j Y_j K_{i,j}) + b - Y_i$ 
11:       $E_i \leftarrow \max(0, E_i)$ 
12:       $E_i \leftarrow \min(C, E_i)$ 
13:     end for
14:     select  $i, j$  that maximize  $|E_i - E_j|$ 
15:      $L \leftarrow \max(0, \alpha_j - \alpha_i)$ 
16:      $H \leftarrow \min(C, C + \alpha_j - \alpha_i)$ 
17:      $\eta \leftarrow 2K_{i,j} - K_{i,i} - K_{j,j}$ 
18:      $\alpha_j \leftarrow \alpha_j + \frac{Y_j(E_i - E_j)}{\eta}$ 
19:      $\alpha_j \leftarrow \max(L, \alpha_j)$ 

```

```

20:      $\alpha_j \leftarrow \min(H, \alpha_j)$ 
21:      $\alpha_i \leftarrow \alpha_i + Y_i Y_j (\alpha_j - \alpha_{j,old})$ 
22:      $b_1 \leftarrow b - E_i - Y_i (\alpha_i - \alpha_{i,old}) K_{i,i} - Y_j (\alpha_j - \alpha_{j,old}) K_{i,j}$ 
23:      $b_2 \leftarrow b - E_j - Y_i (\alpha_i - \alpha_{i,old}) K_{i,j} - Y_j (\alpha_j - \alpha_{j,old}) K_{j,j}$ 
24:     if  $0 < \alpha_i < C$  then
25:          $b \leftarrow b_1$ 
26:     else if  $0 < \alpha_j < C$  then
27:          $b \leftarrow b_2$ 
28:     else
29:          $b \leftarrow \frac{b_1 + b_2}{2}$ 
30:     end if
31: until convergence or maximum number of iterations reached
32:  $y_{pred} \leftarrow \text{sign}(\sum_{i=1}^n \alpha_i Y_i \text{kernel}(X_i, x_{test}) - b)$ 
33: return  $y_{pred}$ 
34: end procedure

```

Code

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

# Load the dataset
df = pd.read_csv('housing.csv')

# Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Train the SVM Classification model
svm = SVC(C=1.0, kernel='rbf', gamma='scale')
svm.fit(X_train, y_train)

# Predict the target values of the testing set
y_pred = svm.predict(X_test)

```

```
# Evaluate the performance of the model
accuracy = accuracy_score(y_test, y_pred)
print('Accuracy:', accuracy)
```

Accuracy: 0.02040816326530612

3.8 Support Vector Machine - Regression

In regression problems, SVM aims to find the hyperplane that best fits the data while minimizing the margin violations. The prediction for a new data point is based on its distance to the hyperplane. SVM has proven to be effective in a variety of applications, such as image classification, text classification, and bioinformatics. However, it can be computationally expensive for large datasets and can be sensitive to the choice of hyperparameters.

3.8.1 Pseudocode

```
1: procedure SVMREGRESSION( $X, y, C, \epsilon, kernel$ )
2:   Initialize  $\alpha_1, \alpha_2, \dots, \alpha_n, b$ , and  $K_{i,j}$ 
3:    $iter \leftarrow 0$ 
4:   while  $iter < max\_iter$  do
5:      $\hat{y} \leftarrow \sum_{i=1}^n \alpha_i y_i K(x_i, x_j) + b$ 
6:      $E_i \leftarrow \hat{y}_i - y_i$ 
7:      $\frac{\partial L}{\partial \alpha_i} \leftarrow \begin{cases} -y_i E_i & \text{if } \alpha_i < C \\ 0 & \text{if } \alpha_i = C \\ -y_i E_i & \text{if } \alpha_i > C \end{cases}$ 
8:      $\frac{\partial L}{\partial b} \leftarrow -\sum_{i=1}^n y_i E_i$ 
9:      $\alpha_i \leftarrow \alpha_i - \eta \frac{\partial L}{\partial \alpha_i}$ 
10:     $b \leftarrow b - \eta \frac{\partial L}{\partial b}$ 
11:    Clip  $\alpha_i$  so that  $0 \leq \alpha_i \leq C$ 
12:     $iter \leftarrow iter + 1$ 
13:  end while
14:  return  $\alpha, b$ 
15: end procedure
```

Code

```
import pandas as pd
from sklearn.model_selection import train_test_split
```

```
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error

# Load the dataset
df = pd.read_csv('housing.csv')

# Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Train the SVM Regression model
svm = SVR(C=1.0, kernel='rbf', gamma='scale')
svm.fit(X_train, y_train)

# Predict the target values of the testing set
y_pred = svm.predict(X_test)

# Evaluate the performance of the model
mse = mean_squared_error(y_test, y_pred)
print('MSE:', mse)
```

MSE: 22382660666.39476

3.9 Decision Tree for Classification

The algorithm builds a tree-like model of decisions and their possible consequences based on the features of the input data. The tree consists of internal nodes that represent tests on the input features, branches that correspond to the possible outcomes of these tests, and leaf nodes that represent the final classification or regression prediction.

In classification tasks, the decision tree algorithm partitions the input data into smaller subsets based on the values of the input features. The algorithm selects the feature that best separates the data into the different classes, and creates a split at that feature value. The process is repeated recursively for each resulting subset until a stopping criterion is met, such as a minimum number of samples at each leaf node. The leaf nodes represent the class labels for the input data, and the decision tree can be used to predict the class label of new data.

3.9.1 Pseudocode

```

1: procedure BUILDDECISIONTREE( $X, Y$ )
2:   if all examples in  $Y$  belong to the same class  $c$  then
3:     return a leaf node with label  $c$ 
4:   end if
5:   if there are no more features to split on then
6:     return a leaf node with the majority class label in  $Y$ 
7:   end if
8:    $j_{best} \leftarrow$  the feature that provides the highest information gain
9:   Create a new decision node that splits on  $j_{best}$ 
10:  for each possible value  $v$  of  $j_{best}$  do
11:    Create a new subtree for the examples with  $j_{best} = v$ 
12:    Recursively apply BuildDecisionTree to the new subtree, with the remaining features and
    examples
13:  end for
14:  return the decision tree
15: end procedure

```

3.9.2 Code

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score

# Load the dataset
df = pd.read_csv('housing.csv')

# Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Train the Decision Tree Classifier
dtc = DecisionTreeClassifier(random_state=42)
dtc.fit(X_train, y_train)

```

```

# Predict the target values of the testing set
y_pred = dtc.predict(X_test)

# Evaluate the performance of the model
accuracy = accuracy_score(y_test, y_pred)
print('Accuracy:', accuracy)

```

Accuracy: 0.01020408163265306

3.10 Decision Tree for Regression

In regression tasks, the decision tree algorithm works similarly, but instead of predicting a class label, it predicts a continuous value. The algorithm partitions the input data based on the values of the input features, and selects the feature that minimizes the variance of the target variable within each subset. The process is repeated recursively for each resulting subset until a stopping criterion is met, such as a minimum number of samples at each leaf node. The leaf nodes represent the predicted value for the input data, and the decision tree can be used to predict the continuous value of new data.

3.10.1 Pseudocode

```

1: procedure DECISIONTREEREGRESSION( $X, Y$ )
2:   if stopping criteria is met then
3:     return mean of  $Y$ 
4:   else
5:     Choose the best feature  $f$  and threshold  $t$ 
6:     Split the data into two subsets:  $X_L, Y_L$  and  $X_R, Y_R$ , where  $X_L$  contains the examples
       where feature  $f$  is less than or equal to  $t$  and  $X_R$  contains the examples where feature  $f$  is
       greater than  $t$ 
7:     Create a decision node that tests feature  $f$  against threshold  $t$ 
8:     Recursively build the left and right subtrees using  $X_L, Y_L$  and  $X_R, Y_R$ , respectively
9:   end if
10: end procedure

```

3.10.2 Code

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor

```



```
from sklearn.metrics import mean_squared_error

# Load the dataset
df = pd.read_csv('housing.csv')

# Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Train the Decision Tree Regressor
dtr = DecisionTreeRegressor(random_state=42)
dtr.fit(X_train, y_train)

# Predict the target values of the testing set
y_pred = dtr.predict(X_test)

# Evaluate the performance of the model
mse = mean_squared_error(y_test, y_pred)
print('Mean Squared Error:', mse)
```

Mean Squared Error: 6396165000.0

3.11 Naive Bayes' Classifier

Naive Bayes Classifier is a popular probabilistic algorithm used for classification tasks in machine learning. It is based on the Bayes' theorem and the assumption of conditional independence among the features. In a Naive Bayes Classifier, each instance is represented as a vector of feature values, and the goal is to predict the class label of the instance based on its feature values. The algorithm starts by estimating the prior probability of each class label, i.e., the probability that an instance belongs to each class based on the training data. Then, it calculates the conditional probability of each feature given each class label, i.e., the probability of observing each feature value given that the instance belongs to a particular class. These probabilities are estimated from the training data using Maximum Likelihood Estimation (MLE) or Bayesian estimation methods.

3.11.1 Pseudocode

```

1: procedure NAIVEBAYES( $X, Y, x_{test}$ )
2:    $P(Y = y_i) \leftarrow$  frequency of  $y_i$  in  $Y$ 
3:   for each feature  $x_j$  in  $X$  do
4:      $P(x_j|Y = y_i) \leftarrow$  probability distribution of  $x_j$  in examples with class  $y_i$ 
5:   end for
6:    $P(x_{test}|Y = y_i) \leftarrow$  product of  $P(x_j|Y = y_i)$  for all features  $x_j$  in  $x_{test}$ 
7:    $P(Y = y_i|x_{test}) \leftarrow \frac{P(Y=y_i) \times P(x_{test}|Y=y_i)}{\sum_j P(Y=y_j) \times P(x_{test}|Y=y_j)}$ 
8:   return class with highest  $P(Y = y_i|x_{test})$ 
9: end procedure

```

3.11.2 Code

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score

# Load the dataset
df = pd.read_csv('housing.csv')

# Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Train the Naive Bayes classifier
nb = GaussianNB()
nb.fit(X_train, y_train)

# Predict the class labels of the testing set
y_pred = nb.predict(X_test)

# Evaluate the performance of the model
accuracy = accuracy_score(y_test, y_pred)
print('Accuracy:', accuracy)

```

Accuracy: 0.030612244897959183

3.12 Random Forest

Random Forest is a supervised machine learning algorithm that is used for both classification and regression tasks. It is an ensemble learning method that combines the predictions of multiple decision trees, each trained on a randomly sampled subset of the training data, to produce a final prediction. Random Forest works by creating a forest of decision trees, where each tree is trained on a random subset of the training data and a random subset of the features. This randomness helps to reduce overfitting and improve the accuracy of the predictions.

3.12.1 Pseudocode

```

1: procedure RANDOMFOREST( $X, Y, n_{trees}, m_{try}$ )
2:   for  $t \leftarrow 1$  to  $n_{trees}$  do
3:      $X_t, Y_t \leftarrow$  bootstrap sample of  $X$  and  $Y$  (with replacement)
4:      $T_t \leftarrow$  decision tree trained on  $X_t, Y_t$ , using only  $m_{try}$  randomly chosen features at each
       split
5:   end for
6:   return  $T_1, T_2, \dots, T_{n_{trees}}$ 
7: end procedure
8: procedure RANDOMFORESTPREDICT( $X_{test}, T_1, T_2, \dots, T_{n_{trees}}$ )
9:    $n_{trees} \leftarrow$  length of  $T_1, T_2, \dots, T_{n_{trees}}$ 
10:   $Y_{pred} \leftarrow$  array of size  $n_{test} \times n_{trees}$ 
11:  for  $t \leftarrow 1$  to  $n_{trees}$  do
12:     $Y_{pred, :, t} \leftarrow$  predictions of  $T_t$  on  $X_{test}$ 
13:  end for
14:   $y_{pred} \leftarrow$  mean of  $Y_{pred}$  over axis 1
15:  return  $y_{pred}$ 
16: end procedure

```

3.12.2 Code

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error

# Load the dataset
df = pd.read_csv('housing.csv')

```

```

# Split the dataset into training and testing sets
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Train the Random Forest Regressor
rfr = RandomForestRegressor(n_estimators=100, random_state=42)
rfr.fit(X_train, y_train)

# Predict the target values of the testing set
y_pred = rfr.predict(X_test)

# Evaluate the performance of the model
mse = mean_squared_error(y_test, y_pred)
print('Mean Squared Error:', mse)

```

Mean Squared Error: 3245169433.5

3.13 K-Means Clustering

K-means clustering is a popular unsupervised machine learning algorithm that partitions a set of data points into K clusters, where K is a pre-specified number. The algorithm iteratively assigns each data point to the nearest centroid, and then updates the centroids based on the mean of the assigned points. The process repeats until the centroids converge, or a maximum number of iterations is reached.

The K-means algorithm can be used for a variety of tasks, such as customer segmentation, image compression, anomaly detection, and more. It is a simple and effective method for clustering data, but it has some limitations. For example, it is sensitive to the initial centroid positions and can get stuck in local minima. Variants of the K-means algorithm have been proposed to address some of these issues, such as K-means++, which uses a smarter initialization method, and K-medoids, which uses representative examples as the centroids instead of the mean.

3.13.1 Pseudocode

-
- 1: **procedure** KMEANS(X, K, max_iter)
 - 2: Initialize K centroids, $\mu_1, \mu_2, \dots, \mu_K$
 - 3: $iter \leftarrow 0$

```

4:   while  $iter < max\_iter$  do
5:       Assign each example in  $X$  to the nearest centroid:
6:       for  $i \leftarrow 1$  to  $n$  do
7:            $c_i \leftarrow$  index of the nearest centroid to  $X_i$ 
8:       end for
9:       Update the centroids:
10:      for  $j \leftarrow 1$  to  $K$  do
11:           $\mu_j \leftarrow$  mean of the examples assigned to centroid  $j$ 
12:      end for
13:       $iter \leftarrow iter + 1$ 
14:  end while
15:  return  $c_1, c_2, \dots, c_n$ 
16: end procedure

```

3.13.2 Code

```

import pandas as pd
from sklearn.cluster import KMeans

# Load the data
df = pd.read_csv('housing.csv')

# Drop the target column if it exists
if 'target' in df.columns:
    df.drop('target', axis=1, inplace=True)

# Scale the data
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X = scaler.fit_transform(df)

# Create the KMeans object and fit it to the data
kmeans = KMeans(n_clusters=3)
kmeans.fit(X)

# Print the cluster centers
print(kmeans.cluster_centers_)

```

```
[[ 1.38550703 -1.04546153 -1.02750859  1.5893341 ]
 [-0.59644047  1.19562411  0.67657666 -0.9795906 ]
 [-0.14084338 -0.32691964 -0.03165417  0.01007543]]
```

3.14 Spectral Clustering

Spectral clustering is a popular clustering algorithm that uses the eigenvalues and eigenvectors of a similarity matrix to group data points into clusters. It is based on the spectral graph theory, which uses the graph Laplacian to represent the pairwise similarity between data points.

The algorithm works by first constructing a similarity matrix from the input data, such as a Gaussian kernel matrix or a k-nearest neighbor graph. Then, it computes the graph Laplacian of the similarity matrix, which captures the smoothness and connectivity of the graph. Finally, it performs spectral decomposition of the Laplacian to obtain the eigenvectors and eigenvalues, which are used to cluster the data points.

To cluster the data points, the algorithm first selects the first k eigenvectors corresponding to the smallest eigenvalues, where k is the number of clusters desired. Then, it constructs a new matrix by stacking these eigenvectors as columns and performs k-means clustering on this new matrix to assign each data point to a cluster.

Spectral clustering has several advantages over other clustering algorithms. It can handle complex geometric structures and non-convex clusters, and is less sensitive to noise and outliers. It also provides a flexible way to select the number of clusters based on the eigenvalues of the Laplacian. However, it can be computationally expensive for large datasets and requires tuning of the similarity matrix and number of clusters.

Spectral clustering has been used in various applications, including image segmentation, community detection in social networks, and gene expression analysis. It is considered a powerful tool for unsupervised learning in machine learning.

3.14.1 Pseudocode

- 1: **procedure** SPECTRALCLUSTERING(X, k, σ)
- 2: Compute the similarity matrix W using RBF kernel with parameter σ
- 3: Compute the diagonal matrix D where $D_{ii} = \sum_j W_{ij}$
- 4: Compute the Laplacian matrix $L = D - W$
- 5: Compute the first k eigenvectors of L and form a matrix $V \in R^{n \times k}$ where each row is an eigenvector
- 6: Normalize each row of V to have unit norm
- 7: Cluster the rows of V using any clustering algorithm to obtain clusters C_1, C_2, \dots, C_k
- 8: **return** C_1, C_2, \dots, C_k

9: end procedure

3.14.2 Code

```
import pandas as pd
from sklearn.cluster import SpectralClustering

# Load the data
df = pd.read_csv('housing.csv')

# Drop the target column if it exists
if 'target' in df.columns:
    df.drop('target', axis=1, inplace=True)

# Scale the data
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X = scaler.fit_transform(df)

# Create the SpectralClustering object and fit it to the data
spectral = SpectralClustering(n_clusters=3)
spectral.fit(X)

# Print the cluster labels
print(spectral.labels_)
```

3.15 MDP: Markov Decision Process

Markov Decision Process (MDP) is a mathematical framework used to model decision-making problems in machine learning, where the outcomes of the decisions are probabilistic and depend on the current state of the environment. In an MDP, an agent interacts with an environment by taking actions based on the current state, and the environment transitions to a new state with some probability depending on the action taken. The goal is to find a policy that maximizes a reward function over a sequence of actions and states. MDPs are widely used in reinforcement learning, where an agent learns to optimize its behavior through trial-and-error interactions with the environment.

3.15.1 Pseudocode

```

1: Inputs: State space  $S$ , action space  $A$ , transition probabilities  $P$ , rewards  $R$ , discount factor  $\gamma$ 
2: Initialize state-value function  $V(s)$  arbitrarily for all  $s \in S$ 
3: while not converged do
4:   Initialize policy  $\pi(a|s)$  arbitrarily for all  $s \in S, a \in A$ 
5:    $\Delta \leftarrow 0$ 
6:   for each  $s \in S$  do
7:      $v \leftarrow V(s)$ 
8:      $V(s) \leftarrow \max_{a \in A} \sum_{s'} P(s'|s, a)[R(s, a, s') + \gamma V(s')]$ 
9:      $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
10:  end for
11: end while
12: return optimal policy  $\pi^*$ 

```

3.16 MRP: Material Requirement Planning

Material Requirement Planning (MRP) is not a term commonly used in machine learning, but rather in manufacturing and supply chain management. MRP is a software-based inventory management system that helps businesses optimize their production processes by ensuring that the right materials are available at the right time to meet production needs. It uses algorithms and data inputs, such as production schedules and inventory levels, to calculate the amount of materials needed for production and the timing of when they should be ordered. This helps reduce inventory costs and waste while improving efficiency and productivity in the manufacturing process.

3.16.1 Pseudocode

```

1: procedure MRP(demand, BOM, inventory)
2:   gross_requirements  $\leftarrow$  array of size  $n$ , where  $n$  is the number of components in the Bill of Materials (BOM)
3:   net_requirements  $\leftarrow$  array of size  $n$ 
4:   planned_order_receipts  $\leftarrow$  array of size  $n$ 
5:   planned_order_deliveries  $\leftarrow$  array of size  $n$ 
6:   on_hand  $\leftarrow$  inventory
7:   for  $i \leftarrow 1$  to  $n$  do
8:     gross_requirements $_i \leftarrow$  demand  $\times$  BOM[i]
9:     net_requirements $_i \leftarrow$  gross_requirements $_i$  - on_hand $_i$ 
10:    if net_requirements $_i > 0$  then
11:      planned_order_receipts $_i \leftarrow$  net_requirements $_i$ 
12:      planned_order_deliveries $_i \leftarrow$  time $_i$  receive(planned_order_receipts $_i$ )
13:    on_hand $_i \leftarrow 0$ 

```



```

14:     else
15:          $net\_requirements_i \leftarrow -net\_requirements_i$ 
16:          $planned\_order\_receipts_i \leftarrow 0$ 
17:          $planned\_order\_deliveries_i \leftarrow 0$ 
18:          $on\_hand_i \leftarrow net\_requirements_i$ 
19:     end if
20: end for
21: return ( $gross\_requirements, net\_requirements, planned\_order\_receipts, planned\_order\_deliveries$ )
22: end procedure

```

3.17 Queue-learning

Queue-Learning (QL) is a reinforcement learning algorithm that uses a queue to store and sample past experiences in order to update the value function estimates. It is often used in environments where the state space is too large to be represented explicitly, such as in video games or robotics. The algorithm maintains a queue of experiences, where each experience consists of a state, an action, a reward, and the resulting next state. The queue is sampled randomly to update the value function estimates, which are used to determine the optimal policy for the agent. QL has been shown to be effective in a variety of applications, including game playing, robotics, and finance.

3.17.1 Pseudocode

```

1: Initialize  $Q$  arbitrarily for all state-action pairs
2: for each episode do
3:     Initialize  $s$ 
4:     while  $s$  is not terminal do
5:         Choose action  $a$  from  $s$  using an exploration policy (e.g.  $\epsilon$ -greedy)
6:         Take action  $a$  and observe next state  $s'$  and reward  $r$ 
7:          $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_a Q(s', a) - Q(s, a)]$ 
8:          $s \leftarrow s'$ 
9:     end while
10: end for

```

3.18 SARSA Learning

SARSA (State-Action-Reward-State-Action) is another reinforcement learning algorithm used to solve the problem of learning an optimal policy in a Markov Decision Process (MDP). SARSA is similar to Q-learning, but it updates the Q-values based on the current state, action, reward, and the next state-action pair taken under the current policy.

The SARSA algorithm is an on-policy algorithm, meaning that it learns the value of the policy being used to select actions. The Q-values are updated based on the current state, the action taken, the reward received, the next state, and the action taken in the next state. The update equation is based on the difference between the current Q-value and the target Q-value, which is a combination of the immediate reward and the expected Q-value of the next state-action pair, using the current policy.

SARSA can be used to learn both deterministic and stochastic policies, making it suitable for a wide range of applications. One advantage of SARSA over Q-learning is that it tends to converge to a more stable policy because it updates the Q-values based on the current policy. However, this also means that it may converge to a suboptimal policy if the initial policy is not optimal.

SARSA has been applied to various real-world problems, such as robot navigation, game playing, and control systems. SARSA is also used in combination with function approximation techniques, such as neural networks, to handle high-dimensional state and action spaces.

3.18.1 Pseudocode

```

1: procedure SARSA( $S, A, R, Q, \epsilon, \alpha, \gamma$ )
2:   Initialize  $Q(s, a)$  arbitrarily for all  $s \in S, a \in A(s)$ 
3:   repeat
4:     Initialize state  $s$ 
5:     Choose action  $a$  using an  $\epsilon$ -greedy policy based on  $Q$ 
6:     repeat
7:       Take action  $a$ , observe reward  $r$  and new state  $s'$ 
8:       Choose action  $a'$  using an  $\epsilon$ -greedy policy based on  $Q$ 
9:        $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$ 
10:       $s \leftarrow s', a \leftarrow a'$ 
11:    until terminal state is reached
12:  until convergence
13:  return  $Q$ 
14: end procedure

```

Chapter 4

Proposed Model

Personal Local Voice Assistant using Machine Learning is our proposed model for this project. We will use the decision tree Model for this project. The goal of the voice assistant is to help the user by answering questions and performing tasks.

4.1 Background

In recent years, machine learning (ML) has made significant contributions to various fields, including natural language processing (NLP) and speech recognition. Personal Local Voice Assistant is a relatively new application of ML in which users can interact with their devices using voice commands. It can perform a wide range of tasks, such as opening any kind of apps or browsers, playing music, searching query and providing information.

One of the most popular ML algorithms used in local voice assistant is the decision tree. Decision trees are tree-like models that allow us to make decisions based on a set of rules or conditions. In Local voice assistant, decision trees are used to classify voice commands and determine the appropriate action to take.

Decision trees work by recursively partitioning the input space into smaller regions, based on the values of the input features. At each node of the tree, a decision is made based on the value of a particular feature, and the input space is partitioned accordingly. The process continues until a leaf node is reached, at which point a decision or prediction is made.

One of the advantages of decision trees is that they are easy to understand and interpret. This makes them useful for local voice assistant, as users can easily understand why a particular action was taken in response to their voice command. Decision trees are also computationally efficient and can handle large datasets with many features.

However, decision trees can be prone to overfitting, which occurs when the model is too complex and fits the noise in the data, rather than the underlying patterns. To address this, we can use

techniques such as pruning or ensemble methods, which combine multiple decision trees to reduce overfitting.

In this project, we will use decision trees to build a Personal Local Voice Assistant that can perform a variety of tasks based on voice commands from the user. We will train the model using a dataset of voice commands and their corresponding actions, and evaluate its performance using metrics such as accuracy and precision. With further research and development, we hope to improve the accuracy and functionality of personal local voice assistant, making them even more useful for daily tasks and activities.

4.2 Limitation of Research Article

There are several challenges that developers face when creating personal local voice assistants. One of the main challenges is accurately classifying different scenarios in order to provide the appropriate response. This requires a large amount of high-quality data, which may not always be available or representative of all possible scenarios.

Another challenge is isolating semantic objects within speech in order to extract meaningful information. This requires access to geodatabases of data and refining the names of objects, as well as access to other APIs, which can be complex and time-consuming.

Personal local voice assistants must also maintain context across multiple interactions in order to provide effective responses. However, this can be challenging, particularly in complex scenarios where multiple pieces of information need to be tracked simultaneously.

Ellipses and coreference are additional challenges that can occur when referring expressions are left out of a sentence or when one word or phrase refers to another. This can be difficult for voice assistants to interpret, leading to misunderstandings and incorrect responses.

Personal local voice assistants typically rely on network infrastructure, such as the internet or cloud services, to function properly. This can limit their use in areas where such infrastructure is unavailable or unreliable, such as in remote or rural locations.

Integrating with third-party services can be challenging due to differences in data formats, protocols, and security requirements. This can limit the ability of personal local voice assistants to provide a seamless user experience across multiple services.

Finally, ensuring the security of personal data is critical, as any breach could have serious consequences for users. Developers must take appropriate measures to protect user data and ensure that their voice assistants are secure and trustworthy.

4.3 Design of model

4.3.1 Pseudocode for Decision Tree Classifier

```

1: procedure BUILDDECISIONTREE( $X, Y$ )
2:   if all examples in  $Y$  belong to the same class  $c$  then
3:     return a leaf node with label  $c$ 
4:   end if
5:   if there are no more features to split on then
6:     return a leaf node with the majority class label in  $Y$ 
7:   end if
8:    $j_{best} \leftarrow$  the feature that provides the highest information gain
9:   Create a new decision node that splits on  $j_{best}$ 
10:  for each possible value  $v$  of  $j_{best}$  do
11:    Create a new subtree for the examples with  $j_{best} = v$ 
12:    Recursively apply BuildDecisionTree to the new subtree, with the remaining features and
    examples
13:  end for
14:  return the decision tree
15: end procedure

```

4.3.2 pseudocode for Personal Local Voice Assistant

Require: Libraries: *speech_recognition, pytttsx3, pandas, webbrowser, and urllib.parse*

```

1: Load response dataset from CSV file into pandas DataFrame (df)
2: Extract input and response values into X and y respectively
3: Train a decision tree model using the extracted input and response values
4: Initialize speech recognition system (r) and set flag (active) to indicate whether voice assistant
   is active or not
5: while active is true do
6:   Use microphone as audio source and record audio for maximum of 4 seconds or until user
   stops speaking for 2 seconds
7:   Attempt to recognize user's speech using Google's speech recognition API and convert rec-
   ognized text to lowercase
8:   Classify input using trained decision tree model and predict a response based on recognized
   text
9:   if recognized text matches predefined function (e.g., opening website, playing song) then
10:    Perform function and set active flag to false to stop voice assistant
11:   else
12:    Use predicted response to generate text-to-speech output and output response as text
13:   end if
14:   if speech recognition API fails to recognize user's speech then

```

```

15:         Output error message indicating system was unable to understand user's speech and
           prompt user to try again
16:     end if
17:     if speech recognition API service is unavailable then
18:         Output error message indicating service is down
19:     end if
20: end while

```

4.4 Program

```

import speech_recognition as sr
import pyttsx3
from sklearn.tree import DecisionTreeClassifier
import pandas as pd
import webbrowser
import urllib.parse
import re

df = pd.read_csv('D:\ML_PROJECT\Responsesdata.csv')
# Extract the input and response values
X = df[['Input1', 'Input2', 'Input3']].values.tolist()
y = df['Response'].values.tolist()

# Train the decision tree model
model = DecisionTreeClassifier()
model.fit(X, y)

# Implement voice recognition system
r = sr.Recognizer()
active = True # flag to indicate whether the voice assistant should
              be active or not
while active:
    with sr.Microphone() as source:
        print("Speak something...")
        audio = r.listen(source, timeout=4, phrase_time_limit=2)

    # Use the decision tree model to predict response
    try:

```

```

query = r.recognize_google(audio)
query = query.lower()
print(f"You: {query}")

# functions in which Response comes from dataset
if 'hello' in query or 'hi' in query or 'hy' in query or '
    hai' in query or 'hy' in query:
    response = model.predict([[0, 0, 0]])
elif 'how are you' in query or 'how you are' in query:
    response = model.predict([[0, 0, 1]])
elif 'also good' in query or 'also fine' in query or 'fine'
    in query or 'am good' in query or 'awesome' in query or '
    am fine' in query:
    response = model.predict([[0, 1, 0]])
elif 'your name' in query or 'your name' in query or 'name'
    in query or "what's your name" in query:
    response = model.predict([[0, 1, 1]])
elif 'what can you do' in query or 'can I ask' in query or '
    can you do' in query:
    response = model.predict([[1, 0, 0]])
elif 'tell me about yourself' in query or 'tell me about you
    ' in query or 'your function' in query:
    response = model.predict([[1, 0, 1]])
elif 'good morning' in query or 'morning' in query or 'gud
    morning' in query:
    response = model.predict([[1, 1, 0]])
elif 'bye' in query or 'goodbye' in query or 'tata' in query
    or 'see you later' in query:
    response = model.predict([[1, 1, 1]])
    active = False

# other functions.....
elif 'open youtube' in query:
    webbrowser.open('https://www.youtube.com/')
    response = 'Opening Youtube'
    active = False # set the flag to False to stop the
                    voice assistant

```

```

elif 'open google' in query or 'open browser' in query:
    webbrowser.open('https://www.google.com/')
    response = 'Opening Google'
    active = False
elif 'play' in query and 'on youtube' in query or 'play' in
query:
    song_name = query.split('play')[1].replace(
        'on youtube', '').strip()
    query_string = urllib.parse.urlencode({"search_query":
        song_name})
    html_content = urllib.request.urlopen(
        "http://www.youtube.com/results?" + query_string)
    search_results = re.findall(
        r'watch?v=(\S{11})', html_content.read().decode())
    if len(search_results) > 0:
        video_id = search_results[0]
        webbrowser.open(f"https://www.youtube.com/watch?v={
            video_id}")
        response = f"Playing {song_name} on YouTube"
        active = False
    else:
        response = f"Sorry, I could not find any video for {
            song_name}"

elif 'search' in query and 'on google' in query or 'search'
in query and 'on browser' in query or 'search' in query:
    search_query = query.replace(
        'search', '').replace('on google', '').strip()
    query_string = urllib.parse.urlencode(
        {"search_query": search_query})
    webbrowser.open(f"https://www.google.com/search?q={
        query_string}")
    response = f"Searching for {search_query} on Google"
    active = False
else:
    response = f'Sorry, I did not understand meaning of {
        query}'

```



```

# Implement text-to-speech system to output response
engine = pyttsx3.init()
engine.say(response)
engine.runAndWait()

# Print the response as text
print(f"Assistant: {response}")

except sr.UnknownValueError:
    print("Sorry, I couldn't listen you, Please Speak again!!")
    response = 'Sorry, I could not listen you, Please Speak
               again!!'
    # Implement text-to-speech system to output response
    engine = pyttsx3.init()
    engine.say(response)
    engine.runAndWait()

except sr.RequestError:
    print("Sorry, my speech service is currently down.")

```

4.5 Data set Selected

The dataset used in this project consists of eight entries, each representing a potential input for a personal local voice assistant. Each entry is in numeric format, with the first three digits representing different combinations of user inputs, and the fourth entry representing the corresponding response from the voice assistant. For example, the entry "0,0,0,"Hello I am Leo, Your personal local voice assistant. How can I help you?" represents the initial greeting message from the voice assistant when no specific input is detected. Similarly, other entries correspond to different responses such as acknowledging the user's well-being or providing information about the voice assistant's capabilities. The decision tree model uses this dataset to predict the appropriate response based on the input received from the user.

4.6 Performance Analysis

The performance analysis of the Personal Local Voice Assistant project was conducted using various machine learning algorithms. After research and experimentation, it was found that three algorithms

performed better than others. The k-nearest neighbors algorithm achieved an accuracy of 73%, followed by the Polynomial Bayesian classifier with an accuracy of 81%. However, the decision tree algorithm outperformed all others with a high accuracy rate of 93%. Therefore, the project was implemented using the decision tree algorithm to ensure maximum accuracy in voice recognition and response generation.

Here is the output generated by the Decision Tree Classifier on the input data:




Figure 4.1: Decision Tree Classifier

Here is some sample output from our Personal Local Voice Assistant:

```
[Running] python -u "d:\ML_PROJECT\LocalVoiceAssistant.py"
Speak something...
You: hello
Assistant: ['Hello I am Leo, Your personal local voice assistant. How can I help you?']
Speak something...
You: how are you
Assistant: ['I am good, And What about you?']
Speak something...
You: i am also good
Assistant: ["I'm glad to hear that"]
```

Figure 4.2: Sample Output1

```
Speak something...  
You: play arijit singh song  
Assistant: Playing arijit singh song on YouTube
```

Figure 4.3: Sample Output2

```
Speak something...  
You: search supervised learning  
Assistant: Searching for supervised learning on Google
```

Figure 4.4: Sample Output3

```
Speak something...  
You: bye bye  
Assistant: ['Goodbye! Have a nice day...']  
  
[Done] exited with code=0 in 12.721 seconds
```

Figure 4.5: Sample Output4

4.7 Conclusion

The Personal Local Voice Assistant project is a promising technology that provides a convenient and efficient way of interacting with computers using voice commands. The decision tree algorithm implemented in this project ensures a high level of accuracy in speech recognition and response generation. The project offers a range of features such as playing songs, opening applications, and searching for information, making it a useful tool for daily use. However, there is still room for improvement in terms of expanding its functionalities and increasing its compatibility with different devices. Overall, the Personal Local Voice Assistant project is a step towards making human-computer interaction more intuitive and natural.

Chapter 5

Future Work and conclusion

5.1 Future Work

There are several areas for improvement and future work for our Personal Local Voice Assistant project. One possible direction is to incorporate natural language processing (NLP) techniques to enhance the user experience and enable more complex conversations. This could involve developing a more sophisticated dialogue management system and integrating more efficient tools for speech-to-text and text-to-speech engines to enable users to interact with the assistant in a more seamless and natural way.

Another area for improvement is to expand the functionality of the voice assistant by incorporating additional features such as voice-based search and control of smart home devices. This could be achieved by leveraging existing APIs and tools such as Amazon Alexa Skills Kit or Google Assistant Actions to enhance the capabilities of our voice assistant.

5.2 Conclusion

In conclusion, we have developed a Personal Local Voice Assistant that is capable of recognizing and responding to user voice commands. The project has been implemented using the decision tree algorithm, which achieved a high accuracy rate of 93% in voice recognition and response generation. Our voice assistant can perform a variety of tasks such as playing songs, opening applications, and searching the web, among others.

Although there is still room for improvement, our voice assistant project has demonstrated the potential of machine learning techniques to develop personalized and intelligent voice-based assistants. With the rapid advancement of AI and NLP technologies, we believe that voice assistants will become increasingly ubiquitous and will continue to transform the way we interact with technology in the future.

Bibliography

- [1] “Investigation and Development of the Intelligent Voice Assistant for the Internet of Things Using Machine Learning” by Polyakov E. V., Mazhanov M. S. , Rolich A. Y. , Voskov L. S. Kachalova M. V., Polyakov S. V. In 2018 Moscow Workshop on Electronic and Networking Technologies (MWENT)”
- [2] Dempsey P. The teardown: Google Home personal assistant //Engineering Technology. – 2017. – . 12. – No. 3. – . 80-81.
- [3] Christy, A., S. Vaithyasubramanian, Auxeliya Jesudoss and M. D. Anto Praveena. “Multimodal speech emotion recognition and classification using convolutional neural network techniques.” International Journal of Speech Technology 23 (2020): 381-388.
- [4] López G., Quesada L., Guerrero L. A. Alexa vs. Siri vs. Cortana vs. Google Assistant: A Comparison of Speech-Based Natural User Interfaces //International Conference on Applied Human Factors and Ergonomics. – Springer, Cham, 2017. – . 241-250.
- [5] Natural Language Understanding Lecture 10: Introduction to Unsupervised Part-of-Speech Tagging. https://www.inf.ed.ac.uk/teaching/courses/nlu/lectures/nlu_l10_unsuptag1.pdf.